

Navigation

Machine Learning Mastery
Making Developers Awesome at Machine Learning

[Click to Take the FREE Crash-Course](#)

Search...



Using CNN for financial time series prediction

by **Adrian Tam** on November 20, 2021 in **Machine Learning for Finance**

34

[Share](#)

[Tweet](#)

[Share](#)

Convolutional neural networks have their roots in image processing. It was first published in LeNet to recognize the MNIST handwritten digits. However, convolutional neural networks are not limited to handling images.

In this tutorial, we are going to look at an example of using CNN for time series prediction with an application from financial markets. By way of this example, we are going to explore some techniques in using Keras for model training as well.

After completing this tutorial, you will know

- What a typical multidimensional financial data series looks like?
- How can CNN applied to time series in a classification problem
- How to use generators to feed data to train a Keras model
- How to provide a custom metric for evaluating a Keras model

Let's get started





Using CNN for financial time series prediction

Photo by [Aron Visuals](#), some rights reserved.

Tutorial overview

This tutorial is divided into 7 parts; they are:

1. Background of the idea
2. Preprocessing of data
3. Data generator
4. The model
5. Training, validation, and test
6. Extensions
7. Does it work?



Background of the idea

In this tutorial we are following the paper titled “CNNpred: CNN-based stock market prediction using a diverse set of variables” by Ehsan Hoseinzade and Saman Haratizadeh. The data file and sample code from the author are available in github:

- <https://github.com/hoseinzadeehsan/CNNpred-Keras>

The goal of the paper is simple: To predict the next day's direction of the stock market (i.e., up or down compared to today), hence it is a binary classification problem. However, it is interesting to see how this problem are formulated and solved.

We have seen the examples on using CNN for sequence prediction. If we consider Dow Jones Industrial Average (DJIA) as an example, we may build a CNN with 1D convolution for prediction. This makes sense because a 1D convolution on a time series is roughly computing its moving average or using digital signal processing terms, applying a filter to the time series. It should provide some clues about the trend.

However, when we look at financial time series, it is quite a common sense that some derived signals are useful for predictions too. For example, price and volume together can provide a better clue. Also some other technical indicators such as the moving average of different window size are useful too. If we put all these align together, we will have a table of data, which each time instance has multiple **features**, and the goal is still to predict the direction of **one** time series.



1	Day	which day of the week	Primitive	Pandas
2	Close	Close price	Primitive	Yahoo Finance
3	Vol	Relative change of volume	Technical Indicator	TA-Lib
4	MOM-1	Return of 2 days before	Technical Indicator	TA-Lib
5	MOM-2	Return of 3 days before	Technical Indicator	TA-Lib
6	MOM-3	Return of 4 days before	Technical Indicator	TA-Lib
7	ROC-5	5 days Rate of Change	Technical Indicator	TA-Lib
8	ROC-10	10 days Rate of Change	Technical Indicator	TA-Lib
9	ROC-15	15 days Rate of Change	Technical Indicator	TA-Lib
10	ROC-20	20 days Rate of Change	Technical Indicator	TA-Lib
11	EMA-10	10 days Exponential Moving Average	Technical Indicator	TA-Lib
12	EMA-20	20 days Exponential Moving Average	Technical Indicator	TA-Lib
13	EMA-50	50 days Exponential Moving Average	Technical Indicator	TA-Lib
14	EMA-200	200 days Exponential Moving Average	Technical Indicator	TA-Lib
15	DTB4WK	4-Week Treasury Bill: Secondary Market Rate	Other	FRED
16	DTB3	3-Month Treasury Bill: Secondary Market Rate	Other	FRED
17	DTB6	6-Month Treasury Bill: Secondary Market Rate	Other	FRED
18	DGS5	5-Year Treasury Constant Maturity Rate	Other	FRED
19	DGS10	10-Year Treasury Constant Maturity Rate	Other	FRED
20	DAAA	Moody's Seasoned Aaa Corporate Bond Yield	Other	FRED
21	DBAA	Moody's Seasoned Baa Corporate Bond Yield	Other	FRED
22	TE1	DGS10-DTB4WK	Other	FRED
23	TE2	DGS10-DTB3	Other	FRED
24	TE3	DGS10-DTB6	Other	FRED
25	TE5	DTB3-DTB4WK	Other	FRED
26	TE6	DTB6-DTB4WK	Other	FRED
27	DE1	DBAA-BAAA	Other	FRED
28	DE2	DBAA-DGS10	Other	FRED
29	DE4	DBAA-DTB6	Other	FRED
30	DE5	DBAA-DTB3	Other	FRED
31	DE6	DBAA-DTB4WK	Other	FRED
32	CTB3M	Change in the market yield on U.S. Treasury securities at 3-month constant maturity, quoted on investment basis	Other	FRED
33	CTB6M	Change in the market yield on U.S. Treasury securities at 6-month constant maturity, quoted on investment basis	Other	FRED
34	CTB1Y	Change in the market yield on U.S. Treasury securities at 1-year constant maturity, quoted on investment basis	Other	FRED
35	Oil	Relative change of oil price (WTI), Oklahoma	Commodity	FRED
36	Oil	Relative change of oil price (Brent)	Commodity	Investing.com
37	Oil	Relative change of oil price (WTI)	Commodity	Investing.com
38	Gold	Relative change of gold price (London market)	Commodity	FRED
39	Gold-F	Relative change of gold price futures	Commodity	Investing.com
40	XAU-USD	Relative change of gold spot U.S. dollar	Commodity	Investing.com
41	XAG-USD	Relative change of silver spot U.S. dollar	Commodity	Investing.com

(continued on next page)

Excerpt from the CNNpred paper showing the list of features used.

Unlike LSTM, which there is an explicit concept of time steps applied, we present data as a matrix in CNN models. As shown in the table below, the features across multiple time steps are presented as a 2D array.

Default															
Options Normal Ruler Zoom															
B2	▲	✖	▼	✖	✖	f _x	10428.049805								
1	Date	A	B	C	D	E	F	G	H	I	J	K	L	M	N
63	3/31/10	10856.6299	-2.9016899	-0.0046565	0.00106091	0.00419341	0.00084404	0.18899232	1.14555378	2.73768116	4.42320615	10823.2468	10724.6936	10552.564	
64	4/1/10	10927.0703	-1.5960622	0.00648824	-0.0046565	0.00106091	0.00419341	0.79198127	1.37209444	2.97055434	4.62393918	10842.1238	10743.9676	10567.2505	
65	4/5/10	10973.5498	-0.7854892	0.00425361	0.00648824	-0.0046565	0.00106091	1.13534896	2.15574155	3.28347802	3.85521382	10866.0194	10765.8326	10583.1838	
66	4/6/10	10969.9902	5.36862338	-0.0003244	0.00425361	0.00648824	-0.0046565	0.68034905	1.7068651	3.08057893	3.95612348	10884.9232	10785.2762	10598.3527	
67	4/7/10	10897.5195	-1.0697488	-0.0066063	-0.0003244	0.00425361	0.00648824	-0.0907675	0.07980153	1.97959432	3.1534236	10887.2135	10795.966	10610.0847	
68	4/8/10	10927.0703	-13.480874	0.0027117	-0.0066063	-0.0003244	0.00425361	0.64882409	0.83904264	1.80181049	3.40426799	10894.4602	10808.4521	10622.5155	
69	4/9/10	10997.3496	-0.6539464	0.00643167	0.0027117	-0.0066063	-0.0003244	0.64316687	1.4402419	2.02408616	3.63282683	10913.1673	10826.4424	10637.2149	



78	4/22/10	11134.29	-1.5243758	0.00084226	0.00070705	0.00225475	0.00666049	-0.0922447	1.89638869	2.55751701	2.75134284	11071.1041	10985.8055	10774.126
79	4/23/10	11204.2803	-2.6734707	0.00628601	0.00084226	0.00070705	0.00225475	1.68459789	1.88164123	2.53691019	3.34898331	11095.318	11006.6126	10790.9948
80	4/26/10	11205.0303	-1.9497073	6.69E-05	0.00628601	0.00084226	0.00070705	1.01857159	1.80865976	2.10944017	3.26873864	11115.2656	11025.5095	10807.2315
81	4/27/10	10991.9902	-6.0532092	-0.0190129	6.69E-05	0.00628601	0.00084226	-1.1250217	-0.2489213	0.20054712	0.88226059	11092.8519	11022.3172	10814.4769
82	4/28/10	11045.2695	-0.7879329	0.0048471	-0.0190129	6.69E-05	0.00628601	-0.7159637	-0.6998116	1.35581312	1.263815	11084.2006	11024.5031	10823.5276
83	4/29/10	11167.3203	-1.6077325	0.01105005	0.0048471	-0.0190129	6.69E-05	0.29665362	0.20413528	2.19866802	2.86175759	11099.3133	11038.1048	10837.0097
84	4/30/10	11008.6104	-4.5531541	-0.014212	0.01105005	0.0048471	-0.0190129	-1.7463855	-0.0912071	0.10239506	0.7462205	11082.8218	11035.2958	10843.7391
85	5/3/10	11151.8301	-0.284403	0.01300979	-0.014212	0.01105005	0.0048471	-0.4747885	0.53894703	1.32528396	1.6246363	11095.3688	11046.3943	10855.8211
86	5/4/10	10926.7695	0.11478752	-0.0201815	0.01300979	-0.014212	0.01105005	-0.5933475	-1.711694	-0.8407919	-0.3939904	11064.7144	11035.0014	10858.6034
87	5/5/10	10866.8301	-0.4213986	-0.0054856	-0.0201815	0.01300979	-0.014212	-1.6155283	-2.3199254	-2.3040343	-0.2816187	11028.7354	11018.9851	10858.926
88	5/6/10	10520.3203	8.83210649	-0.0318869	-0.0054856	-0.0201815	0.01300979	-5.7936907	-5.5142243	-5.6013824	-3.7224067	10936.2963	10971.4932	10845.6474
89	5/7/10	10380.4297	-0.9148126	-0.0132972	-0.0318869	-0.0054856	-0.0201815	-5.7062667	-7.3529987	-5.7922693	-5.6097145	10835.2296	10915.2015	10827.4035
90	5/10/10	10785.1396	-3.7997407	0.03898778	-0.0132972	-0.0318869	-0.0054856	-3.2881637	-3.7473404	-2.7669381	-2.0064573	10826.1224	10902.8146	10825.7461
91	5/11/10	10748.2598	-0.9797658	-0.0034195	0.03898778	-0.0132972	-0.0318869	-1.633692	-2.2173461	-3.3174222	-2.460748	10811.9655	10888.0951	10822.7075
92	5/12/10	10806.0107	10.519840	0.013292019	0.0024165	0.02898779	0.0132922	0.27690237	1.3421029	2.0406407	2.0232055	10927.41	10999.0246	10876.6174

Preprocessing of data

In the following, we try to implement the idea of the CNNpred from scratch using Tensorflow's keras API. While there is a reference implementation from the author in the github link above, we reimplement it differently to illustrate some Keras techniques.

Firstly the data are five CSV files, each for a different market index, under the Dataset directory from github repository above, or we can also get a copy here:

- [CNNpred-data.zip](#)

The input data has a date column and a name column to identify the ticker symbol for the market index. We can leave the date column as time index and remove the name column. The rest are all numerical.

As we are going to predict the market direction, we first try to create the classification label. The market



```
1 ...
2 X["Target"] = (X["Close"].pct_change().shift(-1) > 0).astype(int)
```

The line of code above is to compute the percentage change of the closing index and align the data with the previous day. Then convert the data into either 1 or 0 for whether the percentage change is positive.

For five data file in the directory, we read each of them as a separate pandas DataFrame and keep them in a Python dictionary:

```
1 ...
2 data = {}
3 for filename in os.listdir(DATADIR):
4     if not filename.lower().endswith(".csv"):
5         continue # read only the CSV files
6     filepath = os.path.join(DATADIR, filename)
7     X = pd.read_csv(filepath, index_col="Date", parse_dates=True)
8     # basic preprocessing: get the name, the classification
9     # Save the target variable as a column in dataframe for easier dropna()
10    name = X["Name"][0]
11    del X["Name"]
12    cols = X.columns
13    X["Target"] = (X["Close"].pct_change().shift(-1) > 0).astype(int)
14    X.dropna(inplace=True)
15    # Fit the standard scaler using the training dataset
16    index = X.index[X.index > TRAIN_TEST_CUTOFF]
17    index = index[:int(len(index) * TRAIN_VALID_RATIO)]
18    scaler = StandardScaler().fit(X.loc[index, cols])
19    # Save scale transformed dataframe
20    X[cols] = scaler.transform(X[cols])
21    data[name] = X
```

The result of the above code is a DataFrame for each index, which the classification label is the column "Target" while all other columns are input features. We also normalize the data with a standard scaler.

In time series problems, it is generally reasonable not to split the data into training and test sets randomly, but to set up a cutoff point in which the data before the cutoff is training set while that afterwards is the test set. The scaling above are based on the training set but applied to the entire dataset.



Data generator

We are not going to use all time steps at once, but instead, we use a fixed length of N time steps to predict the market direction at step N+1. In this design, the window of N time steps can start from anywhere. We can just create a large number of DataFrames with large amount of overlaps with one another. To save memory, we are going to build a data generator for training and validation, as follows:

```
1 ...
2 TRAIN_TEST_CUTOFF = '2016-04-21'
3 TRAIN_VALID_RATIO = 0.75
4
5 def datagen(data, seq_len, batch_size, targetcol, kind):
6     "As a generator to produce samples for Keras model"
7     batch = []
8     while True:
9         # Pick one dataframe from the pool
10        key = random.choice(list(data.keys()))
11        df = data[key]
12        input_cols = [c for c in df.columns if c != targetcol]
13        index = df.index[df.index < TRAIN_TEST_CUTOFF]
14        split = int(len(index) * TRAIN_VALID_RATIO)
15        if kind == 'train':
16            index = index[:split] # range for the training set
17        elif kind == 'valid':
18            index = index[split:] # range for the validation set
19        # Pick one position, then clip a sequence length
20        while True:
21            t = random.choice(index) # pick one time step
22            n = (df.index == t).argmax() # find its position in the dataframe
23            if n-seq_len+1 < 0:
24                continue # can't get enough data for one sequence length
25            frame = df.iloc[n-seq_len+1:n+1]
```



```
31     X, y = np.expand_dims(np.array(X), 3), np.array(y)
32     yield X, y
33     batch = []
```

Generator is a special function in Python that does not return a value but to `yield` in iterations, such that a sequence of data are produced from it. For a generator to be used in Keras training, it is expected to `yield` a batch of input data and target. This generator supposed to run indefinitely. Hence the generator function above is created with an infinite loop starts with `while True`.

In each iteration, it randomly pick one DataFrame from the Python dictionary, then within the range of time steps of the training set (i.e., the beginning portion), we start from a random point and take N time steps using the pandas `iloc[start:end]` syntax to create a input under the variable `frame`. This DataFrame will be a 2D array. The target label is that of the last time step. The input data and the label are then appended to the list `batch`. Until we accumulated for one batch's size, we dispatch it from the generator.

The last four lines at the code snippet above is to dispatch a batch for training or validation. We collect the list of input data (each a 2D array) as well as a list of target label into variables `X` and `y`, then convert them into numpy array so it can work with our Keras model. We need to add one more dimension to the numpy array `X` using `np.expand_dims()` because of the design of the network model, as explained below.

The Model

The 2D CNN model presented in the original paper accepts an input tensor of shape $N \times m \times 1$ for N the number of time steps and m the number of features in each time step. The paper assumes $N = 60$ and $m = 82$.

The model comprises of three convolutional layers, as described as follows:



```
5     Input(shape=(seq_len, n_features, 1)),
6     Conv2D(n_filters[0], kernel_size=(1, n_features), activation="relu"),
7     Conv2D(n_filters[1], kernel_size=(3,1), activation="relu"),
8     MaxPool2D(pool_size=(2,1)),
9     Conv2D(n_filters[2], kernel_size=(3,1), activation="relu"),
10    MaxPool2D(pool_size=(2,1)),
11    Flatten(),
12    Dropout(drop_rate),
13    Dense(1, activation="sigmoid")
14  ])
15 return model
```

and the model is presented by the following:

```
1 Model: "sequential"
2 -----
3 Layer (type)          Output Shape       Param #
4 -----
5 conv2d (Conv2D)        (None, 60, 1, 8)    664
6 -----
7 conv2d_1 (Conv2D)      (None, 58, 1, 8)    200
8 -----
9 max_pooling2d (MaxPooling2D) (None, 29, 1, 8) 0
10 -----
11 conv2d_2 (Conv2D)      (None, 27, 1, 8)    200
12 -----
13 max_pooling2d_1 (MaxPooling2D) (None, 13, 1, 8) 0
14 -----
15 flatten (Flatten)     (None, 104)         0
16 -----
17 dropout (Dropout)     (None, 104)         0
18 -----
19 dense (Dense)         (None, 1)           105
20 -----
21 Total params: 1,169
22 Trainable params: 1,169
23 Non-trainable params: 0
```

The first convolutional layer has 8 units, and is applied across all features in each time step. It is followed by a second convolutional layer to consider three consecutive days at once, for it is a common belief that three days can make a trend in the stock market. It is then applied to a max pooling layer and another convolutional layer before it is flattened into a one-dimensional array and applied to a fully-connected layer with sigmoid activation for binary classification.



Training, validation, and test

That's it for the model. The paper used MAE as the loss metric and also monitor for accuracy and F1 score to determine the quality of the model. We should point out that F1 score depends on precision and recall ratios, which are both considering the positive classification. The paper, however, consider the average of the F1 from positive and negative classification. Explicitly, it is the F1-macro metric:

$$F_1 = \frac{1}{2} \left(\frac{\frac{2 \cdot \frac{TP}{TP+FP} \cdot \frac{TP}{TP+FN}}{TP+FP} + \frac{2 \cdot \frac{TN}{TN+FN} \cdot \frac{TN}{TN+FP}}{TN+FN}}{\frac{TP}{TP+FP} + \frac{TN}{TN+FN}} \right)$$

The fraction $\frac{TP}{TP+FP}$ is the precision with TP and FP the number of true positive and false positive.

Similarly $\frac{TP}{TP+FN}$ is the recall. The first term in the big parenthesis above is the normal F1 metric that considered positive classifications. And the second term is the reverse, which considered the negative classifications.

While this metric is available in scikit-learn as `sklearn.metrics.f1_score()` there is no equivalent in Keras. Hence we would create our own by borrowing code from [this stackexchange question](#):

```

1 from tensorflow.keras import backend as K
2
3 def recall_m(y_true, y_pred):
4     true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
5     possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
6     recall = true_positives / (possible_positives + K.epsilon())
7     return recall
8
9 def precision_m(y_true, y_pred):
10    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
11    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
12    precision = true_positives / (predicted_positives + K.epsilon())
13    return precision
14
15 def f1_m(y_true, y_pred):
16    precision = precision_m(y_true, y_pred)
17    recall = recall_m(y_true, y_pred)
18    return 2*((precision*recall)/(precision+recall+K.epsilon()))
19
20 def f1macro(y_true, y_pred):
21    f_pos = f1_m(y_true, y_pred)
22    # negative version of the data and prediction
23    f_neg = f1_m(1-y_true, 1-K.clip(y_pred,0,1))
24    return (f_pos + f_neg)/2

```

The training process can take hours to complete. Hence we want to save the model in the middle of the training so that we may interrupt and resume it. We can make use of checkpoint features in Keras:



```

5     save_best_only=True, save_weights_only=False, save_freq="epoch")
6 ]

```

We set up a filename template `checkpoint_path` and ask Keras to fill in the epoch number as well as validation F1 score into the filename. We save it by monitoring the validation's F1 metric, and this metric is supposed to increase when the model gets better. Hence we pass in the `mode="max"` to it.

It should now be trivial to train our model, as follows:

```

1 seq_len    = 60
2 batch_size = 128
3 n_epochs   = 20
4 n_features = 82
5
6 model = cnnpred_2d(seq_len, n_features)
7 model.compile(optimizer="adam", loss="mae", metrics=["acc", f1macro])
8 model.fit(datagen(data, seq_len, batch_size, "Target", "train"),
9            validation_data=datagen(data, seq_len, batch_size, "Target", "valid"),
10           epochs=n_epochs, steps_per_epoch=400, validation_steps=10, verbose=1,
11           callbacks=callbacks)

```

Two points to note in the above snippets. We supplied "acc" as the accuracy as well as the function `f1macro` defined above as the `metrics` parameter to the `compile()` function. Hence these two metrics will be monitored during training. Because the function is named `f1macro`, we refer to this metric in the checkpoint's `monitor` parameter as `val_f1macro`.

Separately, in the `fit()` function, we provided the input data through the `datagen()` generator as defined above. Calling this function will produce a generator, which during the training loop, batches are fetched from it one after another. Similarly, validation data are also provided by the generator.

Because the nature of a generator is to dispatch data indefinitely. We need to tell the training process on how to define a epoch. Recall that in Keras terms, a batch is one iteration of doing gradient descent update. An epoch is supposed to be one cycle through all data in the dataset. At the end of an epoch is the time to run validation. It is also the opportunity for running the checkpoint we defined above. As Keras has no way to infer the size of the dataset from a generator, we need to tell how many batch it should process in one epoch using the `steps_per_epoch` parameter. Similarly, it is the `validation_steps` parameter to tell how many batch are used in each validation step. The validation does not affect the training, but it will report to us the metrics we are interested. Below is a screenshot of what we will see in the middle of training, which we will see that the metric for training set are updated on each batch but that for validation set is provided only at the end of epoch:

```

1 Epoch 1/20
2 400/400 [=====] - 43s 106ms/step - loss: 0.4062 - acc: 0.6184 - f1ma
3 Epoch 2/20
4 400/400 [=====] - 44s 111ms/step - loss: 0.2760 - acc: 0.7489 - f1ma
5 Epoch 3/20
6 60/400 [----.on+.
7 - ETA: 30s - loss: 0.2300 - acc: 0.7783 - f1macro

```



```
1 def testgen(data, seq_len, targetcol):
2     "Return array of all test samples"
3     batch = []
4     for key, df in data.items():
5         input_cols = [c for c in df.columns if c != targetcol]
6         # find the start of test sample
7         t = df.index[df.index >= TRAIN_TEST_CUTOFF][0]
8         n = (df.index == t).argmax()
9         for i in range(n+1, len(df)+1):
10             frame = df.iloc[i-seq_len:i]
11             batch.append([frame[input_cols].values, frame[targetcol][-1]])
12 X, y = zip(*batch)
13 return np.expand_dims(np.array(X), 3), np.array(y)
14
15 # Prepare test data
16 test_data, test_target = testgen(data, seq_len, "Target")
17
18 # Test the model
19 test_out = model.predict(test_data)
20 test_pred = (test_out > 0.5).astype(int)
21 print("accuracy:", accuracy_score(test_pred, test_target))
22 print("MAE:", mean_absolute_error(test_pred, test_target))
23 print("F1:", f1_score(test_pred, test_target))
```

The structure of the function `testgen()` is resembling that of `datagen()` we defined above. Except in `datagen()` the output data's first dimension is the number of samples in a batch but in `testgen()` is the the entire test samples.

Using the model for prediction will produce a floating point between 0 and 1 as we are using the sigmoid activation function. We will convert this into 0 or 1 by using the threshold at 0.5. Then we use the functions from scikit-learn to compute the accuracy, mean absolute error and F1 score (which accuracy is just one minus the MAE).

Tying all these together, the complete code is as follows:

```
1 import os
2 import random
3
4 import numpy as np
5 import pandas as pd
6 import tensorflow as tf
7 from tensorflow.keras import backend as K
8 from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D, Input
9 from tensorflow.keras.models import Sequential, load_model
10 from tensorflow.keras.callbacks import ModelCheckpoint
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.metrics import accuracy_score, f1_score, mean_absolute_error
13
14 DATADIR = "./Dataset"
15 TRAIN_TEST_CUTOFF = '2016-04-21'
16 TRAIN_VALID_RATIO = 0.75
17
18 # https://datascience.stackexchange.com/questions/45165/how-to-get-accuracy-f1-precision-an
19 # To implement F1 score for validation in a batch
```



```

25
26 def precision_m(y_true, y_pred):
27     true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
28     predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
29     precision = true_positives / (predicted_positives + K.epsilon())
30     return precision
31
32 def f1_m(y_true, y_pred):
33     precision = precision_m(y_true, y_pred)
34     recall = recall_m(y_true, y_pred)
35     return 2*((precision*recall)/(precision+recall+K.epsilon()))
36
37 def f1macro(y_true, y_pred):
38     f_pos = f1_m(y_true, y_pred)
39     # negative version of the data and prediction
40     f_neg = f1_m(1-y_true, 1-K.clip(y_pred,0,1))
41     return (f_pos + f_neg)/2
42
43 def cnnpred_2d(seq_len=60, n_features=82, n_filters=(8,8,8), droprate=0.1):
44     "2D-CNNpred model according to the paper"
45     model = Sequential([
46         Input(shape=(seq_len, n_features, 1)),
47         Conv2D(n_filters[0], kernel_size=(1, n_features), activation="relu"),
48         Conv2D(n_filters[1], kernel_size=(3,1), activation="relu"),
49         MaxPool2D(pool_size=(2,1)),
50         Conv2D(n_filters[2], kernel_size=(3,1), activation="relu"),
51         MaxPool2D(pool_size=(2,1)),
52         Flatten(),
53         Dropout(droprate),
54         Dense(1, activation="sigmoid")
55     ])
56     return model
57
58 def datagen(data, seq_len, batch_size, targetcol, kind):
59     "As a generator to produce samples for Keras model"
60     batch = []
61     while True:
62         # Pick one dataframe from the pool
63         key = random.choice(list(data.keys()))
64         df = data[key]
65         input_cols = [c for c in df.columns if c != targetcol]
66         index = df.index[df.index < TRAIN_TEST_CUTOFF]
67         split = int(len(index) * TRAIN_VALID_RATIO)
68         assert split > seq_len, "Training data too small for sequence length {}".format(seq
69         if kind == 'train':
70             index = index[:split] # range for the training set
71         elif kind == 'valid':
72             index = index[split:] # range for the validation set
73         else:
74             raise NotImplementedError
75         # Pick one position, then clip a sequence length
76         while True:
77             t = random.choice(index) # pick one time step
78             n = (df.index == t).argmax() # find its position in the dataframe
79             if n-seq_len+1 < 0:
80                 continue # this sample is not enough for one sequence length
81             frame = df.iloc[n-seq_len+1:n+1]
82             batch.append([frame[input_cols].values, df.loc[t, targetcol]])

```



```

89         batch = []
90
91     def testgen(data, seq_len, targetcol):
92         "Return array of all test samples"
93         batch = []
94         for key, df in data.items():
95             input_cols = [c for c in df.columns if c != targetcol]
96             # find the start of test sample
97             t = df.index[df.index >= TRAIN_TEST_CUTOFF][0]
98             n = (df.index == t).argmax()
99             # extract sample using a sliding window
100            for i in range(n+1, len(df)+1):
101                frame = df.iloc[i-seq_len:i]
102                batch.append([frame[input_cols].values, frame[targetcol][-1]])
103 X, y = zip(*batch)
104 return np.expand_dims(np.array(X), 3), np.array(y)
105
106 # Read data into pandas DataFrames
107 data = {}
108 for filename in os.listdir(DATADIR):
109     if not filename.lower().endswith(".csv"):
110         continue # read only the CSV files
111     filepath = os.path.join(DATADIR, filename)
112     X = pd.read_csv(filepath, index_col="Date", parse_dates=True)
113     # basic preprocessing: get the name, the classification
114     # Save the target variable as a column in dataframe for easier dropna()
115     name = X["Name"][0]
116     del X["Name"]
117     cols = X.columns
118     X["Target"] = (X["Close"].pct_change().shift(-1) > 0).astype(int)
119     X.dropna(inplace=True)
120     # Fit the standard scaler using the training dataset
121     index = X.index[X.index < TRAIN_TEST_CUTOFF]
122     index = index[:int(len(index) * TRAIN_VALID_RATIO)]
123     scaler = StandardScaler().fit(X.loc[index, cols])
124     # Save scale transformed dataframe
125     X[cols] = scaler.transform(X[cols])
126     data[name] = X
127
128 seq_len = 60
129 batch_size = 128
130 n_epochs = 20
131 n_features = 82
132
133 # Produce CNNpred as a binary classification problem
134 model = cnpred_2d(seq_len, n_features)
135 model.compile(optimizer="adam", loss="mae", metrics=["acc", f1macro])
136 model.summary() # print model structure to console
137
138 # Set up callbacks and fit the model
139 # We use custom validation score f1macro() and hence monitor for "val_f1macro"
140 checkpoint_path = "./cp2d-{epoch}-{val_f1macro:.2f}.h5"
141 callbacks = [
142     ModelCheckpoint(checkpoint_path,
143                     monitor='val_f1macro', mode="max",
144                     verbose=0, save_best_only=True, save_weights_only=False, save_freq="epoch")
145 ]
146 model.fit(datagen(data, seq_len, batch_size, "Target", "train"),

```



```
153 # Test the model
154 test_out = model.predict(test_data)
155 test_pred = (test_out > 0.5).astype(int)
156 print("accuracy:", accuracy_score(test_pred, test_target))
157 print("MAE:", mean_absolute_error(test_pred, test_target))
158 print("F1:", f1_score(test_pred, test_target))
```

Extensions

The original paper called the above model “2D-CNNpred” and there is a version called “3D-CNNpred”. The idea is not only consider the many features of one stock market index but cross compare with many market indices to help prediction on **one index**. Refer to the table of features and time steps above, the data for one market index is presented as 2D array. If we stack up multiple such data from different indices, we constructed a 3D array. While the target label is the same, but allowing us to look at a different market may provide some additional information to help prediction.

Because the shape of the data changed, the convolutional network also defined slightly different, and the data generators need some modification accordingly as well. Below is the complete code of the 3D version, which the change from the previous 2d version should be self-explanatory:

```
1 import os
2 import random
3
4 import numpy as np
5 import pandas as pd
6 import tensorflow as tf
7 from tensorflow.keras import backend as K
8 from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D, Input
9 from tensorflow.keras.models import Sequential, load_model
10 from tensorflow.keras.callbacks import ModelCheckpoint
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.metrics import accuracy_score, f1_score, mean_absolute_error
13
14 DATADIR = "./Dataset"
```



```

21     true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
22     possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
23     recall = true_positives / (possible_positives + K.epsilon())
24     return recall
25
26 def precision_m(y_true, y_pred):
27     true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
28     predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
29     precision = true_positives / (predicted_positives + K.epsilon())
30     return precision
31
32 def f1_m(y_true, y_pred):
33     precision = precision_m(y_true, y_pred)
34     recall = recall_m(y_true, y_pred)
35     return 2*((precision*recall)/(precision+recall+K.epsilon()))
36
37 def f1macro(y_true, y_pred):
38     f_pos = f1_m(y_true, y_pred)
39     # negative version of the data and prediction
40     f_neg = f1_m(1-y_true, 1-K.clip(y_pred,0,1))
41     return (f_pos + f_neg)/2
42
43 def cnnpred_3d(seq_len=60, n_stocks=5, n_features=82, n_filters=(8,8,8), droprate=0.1):
44     "3D-CNNpred model according to the paper"
45     model = Sequential([
46         Input(shape=(n_stocks, seq_len, n_features)),
47         Conv2D(n_filters[0], kernel_size=(1,1), activation="relu", data_format="channels_last"),
48         Conv2D(n_filters[1], kernel_size=(n_stocks,3), activation="relu"),
49         MaxPool2D(pool_size=(1,2)),
50         Conv2D(n_filters[2], kernel_size=(1,3), activation="relu"),
51         MaxPool2D(pool_size=(1,2)),
52         Flatten(),
53         Dropout(droprate),
54         Dense(1, activation="sigmoid")
55     ])
56     return model
57
58 def datagen(data, seq_len, batch_size, target_index, targetcol, kind):
59     "As a generator to produce samples for Keras model"
60     # Learn about the data's features and time axis
61     input_cols = [c for c in data.columns if c[0] != targetcol]
62     tickers = sorted(set(c for _,c in input_cols))
63     n_features = len(input_cols) // len(tickers)
64     index = data.index[data.index < TRAIN_TEST_CUTOFF]
65     split = int(len(index) * TRAIN_VALID_RATIO)
66     assert split > seq_len, "Training data too small for sequence length {}".format(seq_len)
67     if kind == "train":
68         index = index[:split] # range for the training set
69     elif kind == 'valid':
70         index = index[split:] # range for the validation set
71     else:
72         raise NotImplementedError
73     # Infinite loop to generate a batch
74     batch = []
75     while True:
76         # Pick one position, then clip a sequence length
77         while True:
78             t = random.choice(index)

```



```
85     X = np.full(shape, np.nan)
86     for i,ticker in enumerate(tickers):
87         X[i] = frame.xs(ticker, axis=1, level=1).values
88         batch.append([X, data[targetcol][target_index][t]])
89     break
90     # if we get enough for a batch, dispatch
91     if len(batch) == batch_size:
92         X, y = zip(*batch)
93         yield np.array(X), np.array(y)
94         batch = []
95
96 def testgen(data, seq_len, target_index, targetcol):
97     "Return array of all test samples"
98     input_cols = [c for c in data.columns if c[0] != targetcol]
99     tickers = sorted(set(c for _,c in input_cols))
100    n_features = len(input_cols) // len(tickers)
101    t = data.index[data.index >= TRAIN_TEST_CUTOFF][0]
102    n = (data.index == t).argmax()
103    batch = []
104    for i in range(n+1, len(data)+1):
105        # Clip a window of seq_len ends at row position i-1
106        frame = data.iloc[i-seq_len:i]
107        target = frame[targetcol][target_index][-1]
108        frame = frame[input_cols]
109        # convert frame with two level of indices into 3D array
110        shape = (len(tickers), len(frame), n_features)
111        X = np.full(shape, np.nan)
112        for i,ticker in enumerate(tickers):
113            X[i] = frame.xs(ticker, axis=1, level=1).values
114        batch.append([X, target])
115    X, y = zip(*batch)
116    return np.array(X), np.array(y)
117
118 # Read data into pandas DataFrames
119 data = []
120 for filename in os.listdir(DATADIR):
121     if not filename.lower().endswith(".csv"):
122         continue # read only the CSV files
123     filepath = os.path.join(DATADIR, filename)
124     X = pd.read_csv(filepath, index_col="Date", parse_dates=True)
125     # basic preprocessing: get the name, the classification
126     # Save the target variable as a column in dataframe for easier dropna()
127     name = X["Name"][0]
128     del X["Name"]
129     cols = X.columns
130     X["Target"] = (X["Close"].pct_change().shift(-1) > 0).astype(int)
131     X.dropna(inplace=True)
132     # Fit the standard scaler using the training dataset
133     index = X.index[X.index < TRAIN_TEST_CUTOFF]
134     index = index[:int(len(index) * TRAIN_VALID_RATIO)]
135     scaler = StandardScaler().fit(X.loc[index, cols])
136     # Save scale transformed dataframe
137     X[cols] = scaler.transform(X[cols])
138     data[name] = X
139
140 # Transform data into 3D dataframe (multilevel columns)
141 for key, df in data.items():
142     df.columns = pd.MultiIndex.from_product([df.columns, [key]])
```



```

149 n_stocks = 5
150
151 # Produce CNNpred as a binary classification problem
152 model = cnnpred_3d(seq_len, n_stocks, n_features)
153 model.compile(optimizer="adam", loss="mae", metrics=["acc", "f1macro])
154 model.summary() # print model structure to console
155
156 # Set up callbacks and fit the model
157 # We use custom validation score f1macro() and hence monitor for "val_f1macro"
158 checkpoint_path = "./cp3d-{epoch}-{val_f1macro:.2f}.h5"
159 callbacks = [
160     ModelCheckpoint(checkpoint_path,
161                     monitor='val_f1macro', mode="max",
162                     verbose=0, save_best_only=True, save_weights_only=False, save_freq="epoch")
163 ]
164
165 model.fit(datagen(data, seq_len, batch_size, "DJI", "Target", "train"),
166             validation_data=datagen(data, seq_len, batch_size, "DJI", "Target", "valid"),
167             epochs=n_epochs, steps_per_epoch=400, validation_steps=10, verbose=1, callbacks=callbacks)
168
169 # Prepare test data
170 test_data, test_target = testgen(data, seq_len, "DJI", "Target")
171
172 # Test the model
173 test_out = model.predict(test_data)
174 test_pred = (test_out > 0.5).astype(int)
175 print("accuracy:", accuracy_score(test_pred, test_target))
176 print("MAE:", mean_absolute_error(test_pred, test_target))
177 print("F1:", f1_score(test_pred, test_target))

```

While the model above is for next-step prediction, it does not stop you from making prediction for k steps ahead if you replace the target label to a different calculation. This may be an exercise for you.

Does it work?

As in all prediction projects in the financial market, it is always unrealistic to expect a high accuracy. The training parameter in the code above can produce slightly more than 50% accuracy in the testing set. While the number of epochs and batch size are deliberately set smaller to save time, there should not be much room for improvement.

In the original paper, it is reported that the 3D-CNNpred performed better than 2D-CNNpred but only attaining the F1 score of less than 0.6. This is already doing better than three baseline models mentioned in the paper. It may be of some use, but not a magic that can help you make money quick.

From machine learning technique perspective, here we classify a panel of data into whether the market



Further readings

The original paper is available at:

- “CNNPred: CNN-based stock market prediction using several data sources”, by Ehsan Hoseinzadeh, Saman Haratizadeh, 2019.
(<https://arxiv.org/abs/1810.08923>)

If you are new to finance application and want to build the connection between machine learning techniques and finance, you may find this book useful:

- *Machine Learning in Finance: From Theory to Practice*, by Matthew F. Dixon, Igor Halperin, and Paul Bilokon. 2000.
(<https://www.amazon.com/dp/3030410676/>)

On the similar topic, we have a previous post on using CNN for time series, but using 1D convolutional layers;

- [How to develop convolutional neural network models for time series forecasting](#)

You may also find the following documentation helpful to explain some syntax we used above:

- Pandas user guide: https://pandas.pydata.org/pandas-docs/stable/user_guide/index.html
- Keras model training API: https://keras.io/api/models/model_training_apis/
- Keras callbacks API: <https://keras.io/api/callbacks/>



Summary

In this tutorial, you discovered how a CNN model can be built for prediction in financial time series.

Specifically, you learned:

- How to create 2D convolutional layers to process the time series
- How to present the time series data in a multidimensional array so that the convolutional layers can be applied
- What is a data generator for Keras model training and how to use it
- How to monitor the performance of model training with a custom metric
- What to expect in predicting financial market

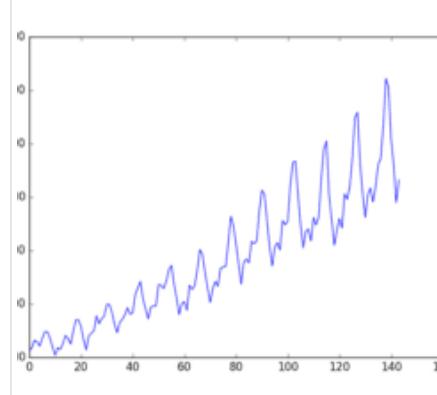
[Share](#)

Tweet

[Share](#)

More On This Topic

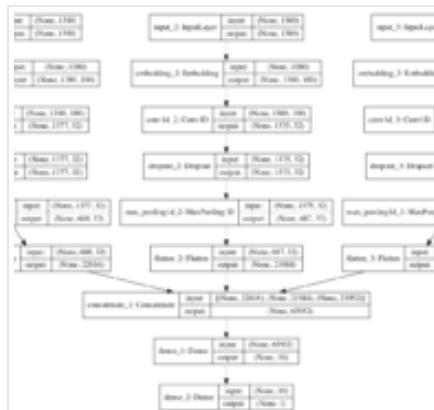
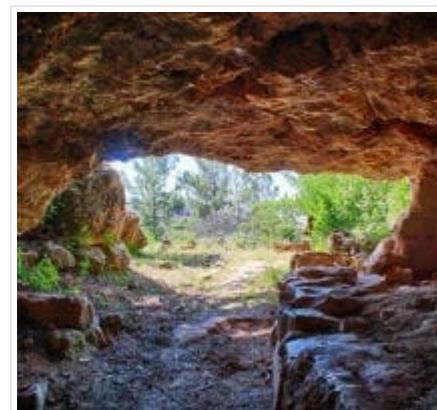




[Understand Time Series Forecast Uncertainty Using...](#)

[Time Series Prediction with Deep Learning in Keras](#)

[Time Series Prediction with LSTM Recurrent Neural...](#)



[LSTM for Time Series Prediction in PyTorch](#)

[CNN Long Short-Term Memory Networks](#)

[How to Develop a Multichannel CNN Model for Text...](#)



About Adrian Tam

Adrian Tam, PhD is a data scientist and software engineer.

[View all posts by Adrian Tam →](#)

< Face Recognition using Principal Component Analysis

Visualizing the vanishing gradient problem >



REPLY ↩



Paul Kornreich November 19, 2021 at 2:40 am #

This has promise,, especially using multiple parameters, but, in general, CNNs are losing out in accuracy compared with Transformers.



Adrian Tam November 19, 2021 at 10:31 am #

REPLY ↗

Correct. But this shows how simple it is to get something not too far away.



tanunchai November 19, 2021 at 5:56 am #

REPLY ↗

There is error at a line in def datagen() on the following:

```
index = data.index[data.index < TRAIN_TEST_CUTOFF]
```

compliler said

"File "C:\Users\TANUNC~1.J\AppData\Local\Temp\ipykernel_10252/1129467454.py", line 66

```
index = data.index[data.index < TRAIN_TEST_CUTOFF]
```

^

SyntaxError: invalid syntax

"

How to solve it ?

waiting your answer, Thanks in advance



Adrian Tam November 19, 2021 at 10:39 am #

REPLY ↗

The line seems OK but maybe the line before it in your code caused some problem. Can you check if you copied the code correctly?





William Smith November 19, 2021 at 9:44 am #

REPLY ↗

@tanunchai

The code above seems to say

```
index = data.index[data.index < TRAIN_TEST_CUTOFF]
```

Try changing to

```
index = data.index[data.index < TRAIN_TEST_CUTOFF]
```



William Smith November 19, 2021 at 9:45 am #

REPLY ↗

Argh. The code actually says [data.index & I t ; TRAIN_TEST_CUTOFF]

Change to

```
index = data.index[data.index < TRAIN_TEST_CUTOFF]
```



Adrian Tam November 19, 2021 at 10:57 am #

REPLY ↗

Thanks William. There were some hassles with the HTML. I fixed that and the copy should work now.



tanunchai November 19, 2021 at 10:09 pm #

REPLY ↗

Thanks you William Smith



tanunchai November 19, 2021 at 10:17 pm #

REPLY ↗

How to solve this bug ?

Now facing new problem, in def datagen() at nest loop while true has the problem on the following:

```
if n-seq_len+1 & 0:
```

```
while True:
```



```
if n-seq_len+1 & 0: # ****error said invalid syntax  
continue # this sample is not enough for one sequence length  
frame = data.iloc[n-seq_len+1:n+1][input_cols]  
# convert frame with two l
```



tanunchai November 19, 2021 at 10:20 pm <#>

REPLY ↗

Now facing new problem, in def datagen() at nest loop while true has the problem on the following:

```
if n-seq_len+1 < 0:  
  
while True:  
# Pick one position, then clip a sequence length  
while True:  
t = random.choice(index)  
n = (data.index == t).argmax()  
#index = data.index[data.index < TRAIN_TEST_CUTOFF] , the line below has the problem  
if n-seq_len+1 < 0:  
continue # this sample is not enough for one sequence length  
frame = data.iloc[n-seq_len+1:n+1][input_cols]  
# convert frame with two level of indices into 3D array  
shape = (len(tickers), len(frame), n_features)  
X = np.full(shape, np.nan)  
for i,ticker in enumerate(tickers):
```



tanunchai November 19, 2021 at 10:26 pm <#>

REPLY ↗

at def datagen() still has problem on the following;

```
if n-seq_len+1 < 0: # Error line , say "invalid syntax"  
&lt doing what ? I do not understand and it made error also.
```

```
while True:  
# Pick one position, then clip a sequence length  
while True:  
t = random.choice(index)  
n = (data.index == t).argmax()
```



```
# convert frame with two level of indices into 3D array  
shape = (len(tickers), len(frame), n_features)
```



tanunchai November 19, 2021 at 10:34 pm #

REPLY ↗

ModuleNotFoundError: No module named 'f1metrics'

How to solve this bug ?

ModuleNotFoundError Traceback (most recent call last)

```
C:\Users\TANUNC~1.J\AppData\Local\Temp\ipykernel_2580/1760727365.py in  
12 from sklearn.metrics import accuracy_score, f1_score, mean_absolute_error  
13  
—> 14 from f1metrics import f1macro  
15  
16 DATADIR = "./Dataset"
```

ModuleNotFoundError: No module named 'f1metrics'



Adrian Tam November 20, 2021 at 2:33 am #

REPLY ↗

Sorry for all these hassles. Something wrong with the plugin for rendering the code here caused all the mess. I fixed it for now, so please copy over the code and try again.



Jack November 20, 2021 at 5:39 am #

REPLY ↗

Interesting! What is the difference between using:

- conv1d with kernel_size=n_features and input size N x m
- conv2d with kernel_size=(1, n_features) and input size N x m x 1

Are these two equivalent?



Adrian Tam November 20, 2021 at 1:36 pm #

REPLY ↗



**Doni** November 21, 2021 at 2:52 pm #[REPLY ↗](#)

how to deploy this code on the web app. ?

**Adrian Tam** November 23, 2021 at 1:06 pm #[REPLY ↗](#)

That's a vague question – you need to think about what the web app expects and how to wrap the model into a function to talk to the web app

**kaplan** November 29, 2021 at 7:36 am #[REPLY ↗](#)

Is FFDNet (fast and flexible denoising convolutional neural networks) suitable for financial time series?

**Adrian Tam** November 29, 2021 at 8:56 am #[REPLY ↗](#)

Haven't tried that.

**Priya** December 1, 2021 at 9:14 pm #[REPLY ↗](#)

My question is from the blog <https://machinelearningmastery.com/feature-selection-machine-learning-python/>

Sorry for asking this question here.

My work is to develop a model for multioutput prediction (i.e., predicting five outputs by a single model). When I applied Kbest and recursive feature elimination methods to select the best features, I got an error 'bad input shape (x, 5)' (5 is output vectors here). However, PCA works well as it doesn't depend upon the output vector.

Does it mean that we can apply these feature selection algorithms (Kbest and RFE) only for a single output prediction problem?

**Adrian Tam** December 2, 2021 at 2:55 am #[REPLY ↗](#)[REPLY ↗](#)

**Pete** March 28, 2022 at 4:32 am #

I think there is too much information in the csv file. Instead of looking too much data and try to make a fundamental analysis it can also be tried to find a 'pattern'. With pattern I mean looking only to the open price, close price, max price and min price information of each candlestick.

For example:

Let's consider the candlestick chart of D1 (1 candlestick = 1 day info). Objective: predict the fifth candlestick (up or down) with the info of the 4 formed before. The less quantity of candlestick the easier a pattern can be found

I will try modifying the input data with this model and let you know the results

**James Carmichael** March 28, 2022 at 7:05 am #[REPLY ↩](#)

Thanks for the feedback Pete! Please share your findings.

**Bojie** May 10, 2022 at 8:28 am #[REPLY ↩](#)

Many thanks for your tutorial, which is very helpful.

Can you please let me know how to write a data generator to read all different classes of .txt files (likes the function of "datagen.flow_from_directory" which read all different classes of image files)?

**Yu** June 15, 2022 at 4:10 am #[REPLY ↩](#)

Thanks for this useful article!

I was trying to run your codes a couple of times but there are certain instability issues happened. Sometimes the prediction rate is about 51%, which is good result, but there are also times the prediction rate will go down to 48%. My guess is that this is related to the initialization step. Do you have any insights on this issue?

**James Carmichael** June 15, 2022 at 7:16 am #[REPLY ↩](#)

Hi Yu...The following resource should add clarity:

[DEDIV ↩](#)

**Yu** June 16, 2022 at 7:10 am #

Hi, James,

Thanks for your suggestion! I do consider the multi-restart idea and similarly using ensemble of results. But I didn't have any luck for that yet either.

Do you think for a larger dataset this issue will be alleviated?

**rk** November 26, 2022 at 5:55 am #**REPLY ↗**

Hi, Could you help me i am having an issue with the code

InvalidArgumentError: Graph execution error

It is associated with [[model.fit(datagen(data, seq_len, batch_size,"Target","train"), validation_data=datagen(data, seq_len, batch_size, "Target", "valid"), epochs=n_epochs, steps_per_epoch=400, validation_steps=10, verbose=1, callbacks=callbacks)]]

I was able to trace it back to this line within the code in datagen (2DConv):

```
input_cols = [c for c in df.columns if c != targetcol]
```

Im using google colab and it says that when i try to single it out and run it

ValueError: The truth value of a Series is ambiguous. Use a.empty, a.bool(), a.item(), a.any() or a.all().

After trying: input_cols = [c for c in df.columns if c != targetcol.all()]

The issue persists, i would appreciate your help on it

**James Carmichael** November 26, 2022 at 8:22 am #**REPLY ↗**

Hi rk...did you copy and paste the code or type it in? Also, you may want to try to execute your code in Google Colab.

**de santos** March 25, 2023 at 8:03 pm #**REPLY ↗**

Hello, Can you tell me what the following lines of code do? They are:

```
name = X["Name"][0]
```



Since I get a bug here on google colab that return the error :

KeyError : "name"

This only happen when I use my own .csv data files. When I ran the codes using your dataset .csv files, I did not encounter any problem though. My csv files only have 7 columns : Date,Close,Open, Hlgh, Low, Vol, Change %.

From what I understood, you are trying to drop the other columns but the Date and Close columns, right?

Thanks in advance!



James Carmichael March 26, 2023 at 10:33 am #

REPLY ↗

Hi de santos... You are correct. Can you provide any more detail about the error? Is that the entire error message?



de santos June 27, 2023 at 6:57 pm #

REPLY ↗

Thanks for replying, James! I thought my response was not approved,hence the late reply 4 months later, I'm so sorry. There was another error about Integer data type since originally the Vol column had a word "k" as substitute for "thousand", that can be easily fixed by Find/Replace function of Excel to uniform the data into integer numbers. Other than that, the "Name" error still persists, here is the full error message:

KeyError Traceback (most recent call last)

```
/usr/local/lib/python3.10/dist-packages/pandas/core/indexes/base.py in get_loc(self, key, method, tolerance)
3801 try:
-> 3802     return self._engine.get_loc(casted_key)
3803 except KeyError as err:
```

4 frames

```
pandas/_libs/hashtable_class_helper.pxi in
pandas._libs.hashtable.PyObjectHashTable.get_item()
```

```
pandas/_libs/hashtable_class_helper.pxi in
pandas._libs.hashtable.PyObjectHashTable.get_item()
```

KeyError: 'Name'

The above exception was the direct cause of the following exception:



3803 except KeyError as err:
-> 3804 raise KeyError(key) from err
3805 except TypeError:
3806 # If we have a listlike key, _check_indexing_error will raise

KeyError: 'Name'



the broad, jeff October 16, 2023 at 2:38 am #

REPLY ↲

I am confused with the below code and would like to seek some clarification:

```
while True:  
    t = random.choice(index) # pick one time step
```

Won't this line of code, always pick up from an index randomly, thus breaking sequence of time series data. As I understand in time series problems, the sequence of time series should be preserved?

What am I missing here? Can you please help me understand this better?



James Carmichael October 16, 2023 at 6:05 am #

REPLY ↶

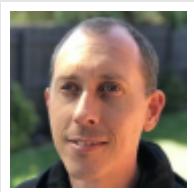
Hi...This line simply picks a random starting point in the sequence. From there, the sequence is preserved.

[Leave a Reply](#)

Name (required)



SUBMIT COMMENT



Welcome!

I'm *Jason Brownlee* PhD
and I **help developers** get results with **machine learning**.
[Read more](#)

Never miss a tutorial:



Picked for you:



[Your First Deep Learning Project in Python with Keras Step-by-Step](#)



[Your First Machine Learning Project in Python Step-By-Step](#)



[How to Develop LSTM Models for Time Series Forecasting](#)





Machine Learning for Developers

Loving the Tutorials?

The [EBook Catalog](#) is where
you'll find the **Really Good** stuff.

[>> SEE WHAT'S INSIDE](#)

© 2024 Guiding Tech Media. All Rights Reserved.

[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

